

stat602_hw2

4.3

First we read in the data set and remove the irrelevant columns.

```
house <- read_excel("~/work/homework/AmesHousingData.xlsx")
house <- data.frame(Price = house$Price, Size = house$Size,
                    Fireplace = house$Fireplace,
                    Basementbath = house`Bsmt Bath`, Land = house$Land, intersect = 1)
head(house)

##      Price  Size Fireplace Basementbath   Land intersect
## 1 141900 1762          0         1 7200        1
## 2 124000 2104          0         0 6300        1
## 3 109900 1040          0         0 6240        1
## 4 105000 1220          0         0 6120        1
## 5  84900 1177          0         0 7218        1
## 6 200000 2234          2         1 12975        1
```

Now we need a matrix of all the possible regressor combinations.

```
regMat <- expand.grid(c(TRUE, FALSE), c(TRUE, FALSE),
                       c(TRUE, FALSE), c(TRUE, FALSE))
regMat <- cbind(TRUE, regMat)

head(regMat)

##      TRUE Var1  Var2  Var3 Var4
## 1  TRUE  TRUE  TRUE  TRUE TRUE
## 2  TRUE FALSE  TRUE  TRUE TRUE
## 3  TRUE  TRUE FALSE  TRUE TRUE
## 4  TRUE FALSE FALSE  TRUE TRUE
## 5  TRUE  TRUE  TRUE FALSE TRUE
## 6  TRUE FALSE  TRUE FALSE TRUE
```

Then we construct a formula so that from each row of regMat we can determine the the regressor combinations.

```
formu <- function(vec)
{
  vec <- as.matrix(vec)
  regressors <- c("intersect", "Size", "Fireplace", "Basementbath", "Land")
  out <- as.formula(paste(c("Price ~ 0", regressors[vec]), collapse=" + "))
```

```

    return(out)
}

formu(regMat[1,])

## Price ~ 0 + intersect + Size + Fireplace + Basementbath + Land
## <environment: 0xcb794c8>

```

Cross-validation using caret package

We will repeat the 8-fold crossval times.

```

ctrl <- trainControl(method = "repeatedcv", repeats = 5, number = 8)

RMSE <- numeric(16)
for(i in 1:length(RMSE)){
  model_caret <- train(formu(regMat[i,]), data = house, trControl = ctrl, method = "lm")
  RMSE[i] = model_caret$results[,2]
}
RMSE

## [1] 21709.31 26670.86 24488.29 32482.33 22305.24 27720.10 25184.11
## [8] 34037.32 22336.93 27915.77 26790.50 36653.70 23441.72 28786.40
## [15] 27497.23 36665.69

order(RMSE)

## [1] 1 5 9 13 3 7 2 11 15 6 10 14 4 8 12 16

regMat

##      TRUE Var1 Var2 Var3 Var4
## 1   TRUE  TRUE  TRUE  TRUE
## 2   TRUE FALSE  TRUE  TRUE
## 3   TRUE  TRUE FALSE  TRUE
## 4   TRUE FALSE FALSE  TRUE
## 5   TRUE  TRUE FALSE  TRUE
## 6   TRUE FALSE  TRUE FALSE
## 7   TRUE  TRUE FALSE FALSE
## 8   TRUE FALSE FALSE FALSE
## 9   TRUE  TRUE  TRUE FALSE
## 10  TRUE FALSE  TRUE FALSE
## 11  TRUE  TRUE FALSE FALSE
## 12  TRUE FALSE FALSE FALSE
## 13  TRUE  TRUE  TRUE FALSE
## 14  TRUE FALSE  TRUE FALSE
## 15  TRUE  TRUE FALSE FALSE
## 16  TRUE FALSE FALSE FALSE

```

We can see that the full model has the smallest mean squared prediction error.

4.4

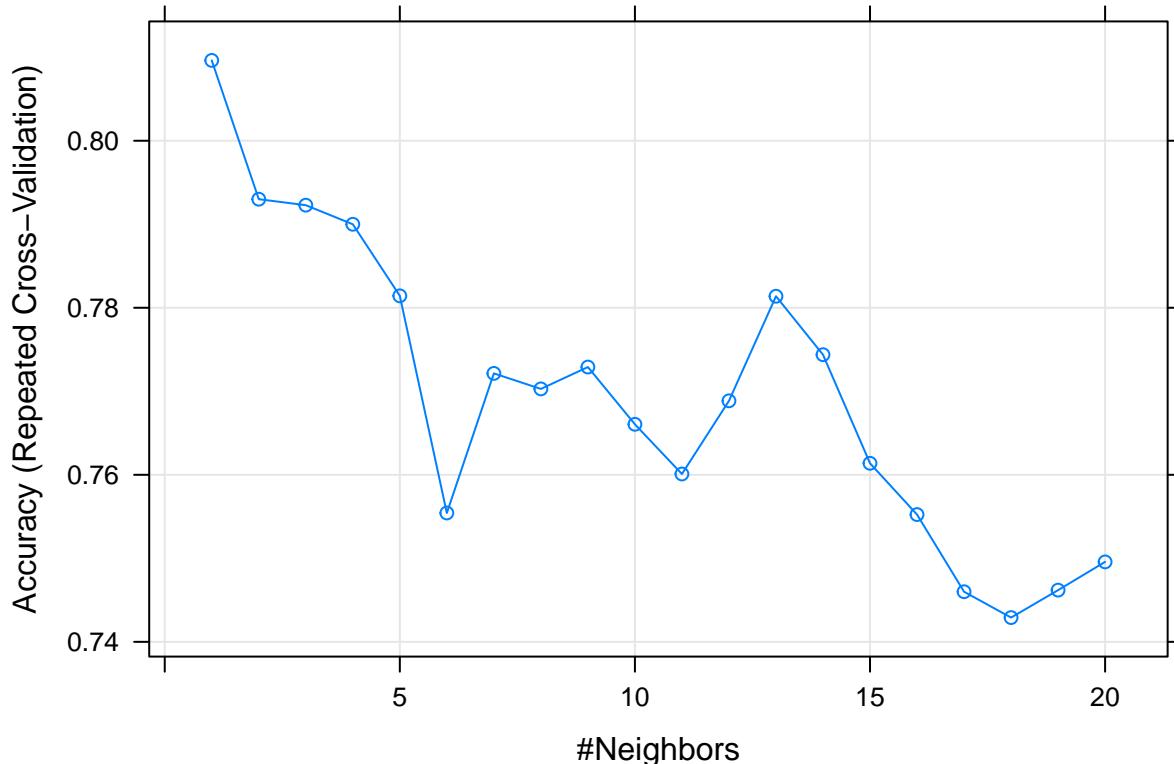
First read in the dataset:

```
glass <- read.csv("~/work/homework/glass.data", header=FALSE)
names(glass) <- c("Id", "RI", "Na", "Mg", "Al", "Si", "K", "Ca", "Ba", "Fe", "Type")
glass_sub <- subset(glass, Type %in% c(1,2))
glass_sub$Type <- as.factor(glass_sub$Type)
```

(a)

The best number of neighbors for this prediction task is 1.

```
ctrl <- trainControl(method="repeatedcv", repeats = 10, number = 10)
knnFit <- train(Type ~ . - Id, data = glass_sub, method = "knn", trControl = ctrl, preProcess = c("center", "scale"))
plot(knnFit)
```



(b)

See hand-written solution.

5.1

(a)

```
X <- c(2, 4, 3, 5, 1, 4, 3, 4, 2, 3, 7, 5, 6, 4, 4, 2, 5, 1, 2, 4)
X <- matrix(X, nrow = 5)

qrObj <- qr(X)
svdObj <- svd(X)
```

The Q matrix and R matrix are:

```
qr.Q(qrObj)
```

```
##          [,1]      [,2]      [,3]      [,4]
## [1,] -0.2696799  0.57022480  0.7723433  0.04477265
## [2,] -0.5393599 -0.06579517 -0.1252449  0.56476381
## [3,] -0.4045199  0.37283929 -0.3548604 -0.70772210
## [4,] -0.6741999 -0.50442963  0.1043707 -0.12567762
## [5,] -0.1348400  0.52636136 -0.5009794  0.40295387
```

```
qr.R(qrObj)
```

```
##          [,1]      [,2]      [,3]      [,4]
## [1,] -7.416198 -6.067799 -10.247838 -5.528439
## [2,]  0.000000  4.145096  5.987360  2.280899
## [3,]  0.000000  0.000000  1.064581 -1.231574
## [4,]  0.000000  0.000000  0.000000  3.566102
```

The columns of Q matrix give the bases for $C(\mathbf{X})$.

The SVD results are:

```
svdObj
```

```
## $d
## [1] 16.5812586  3.7842977  3.3820909  0.5499212
##
## $u
##          [,1]      [,2]      [,3]      [,4]
## [1,] -0.4993567  0.5283913  0.1647824 -0.66384835
## [2,] -0.5040186 -0.5891454  0.1259937 -0.11349453
## [3,] -0.4582533  0.4878118 -0.2526166  0.64627429
## [4,] -0.3913976 -0.3254921 -0.6846703 -0.08815998
## [5,] -0.3652670 -0.1726408  0.6514474  0.34782427
##
## $v
##          [,1]      [,2]      [,3]      [,4]
## [1,] -0.4047825 -0.4324356 -0.79720334  0.11669381
## [2,] -0.4354972  0.2981978  0.18084854  0.82988796
## [3,] -0.7111504  0.4458822  0.03987313 -0.54209258
## [4,] -0.3751780 -0.7247528  0.57460488 -0.06167786
```

The U has columns spanning $C(\mathbf{X})$.

(b)

The eigen decomposition of $X'X$ is:

```
(V <- svdObj$v)
```

```
## [,1]      [,2]      [,3]      [,4]
## [1,] -0.4047825 -0.4324356 -0.79720334 0.11669381
## [2,] -0.4354972  0.2981978  0.18084854 0.82988796
## [3,] -0.7111504  0.4458822  0.03987313 -0.54209258
## [4,] -0.3751780  -0.7247528  0.57460488 -0.06167786
```

```
(D_sq <- diag(svdObj$d) %*% diag(svdObj$d))
```

```
## [,1]      [,2]      [,3]      [,4]
## [1,] 274.9381 0.00000 0.00000 0.0000000
## [2,] 0.0000 14.32091 0.00000 0.0000000
## [3,] 0.0000 0.00000 11.43854 0.0000000
## [4,] 0.0000 0.00000 0.00000 0.3024134
```

The eigen decomposition of XX' is:

```
(U <- svdObj$u)
```

```
## [,1]      [,2]      [,3]      [,4]
## [1,] -0.4993567 0.5283913 0.1647824 -0.66384835
## [2,] -0.5040186 -0.5891454 0.1259937 -0.11349453
## [3,] -0.4582533 0.4878118 -0.2526166 0.64627429
## [4,] -0.3913976 -0.3254921 -0.6846703 -0.08815998
## [5,] -0.3652670 -0.1726408 0.6514474 0.34782427
```

```
(D_sq <- diag(svdObj$d) %*% diag(svdObj$d))
```

```
## [,1]      [,2]      [,3]      [,4]
## [1,] 274.9381 0.00000 0.00000 0.0000000
## [2,] 0.0000 14.32091 0.00000 0.0000000
## [3,] 0.0000 0.00000 11.43854 0.0000000
## [4,] 0.0000 0.00000 0.00000 0.3024134
```

(c)

rank = 1

```
( rank1approx <- best_approx(X, 1) )
```

```
## [,1]      [,2]      [,3]      [,4]
## [1,] 3.351583 3.605900 5.888298 3.106459
## [2,] 3.382873 3.639565 5.943271 3.135461
## [3,] 3.075706 3.309089 5.403616 2.850758
## [4,] 2.626983 2.826318 4.615269 2.434854
## [5,] 2.451600 2.637627 4.307144 2.272298
```

```
rank = 2
```

```
( rank2approx <- best_approx(X, 2) )
```

```
## [,1]      [,2]      [,3]      [,4]
## [1,] 2.486889 4.202174 6.779880 1.657251
## [2,] 4.346989 2.974732 4.949176 4.751299
## [3,] 2.277419 3.859570 6.226726 1.512847
## [4,] 3.159640 2.459010 4.066050 3.327575
## [5,] 2.734121 2.442807 4.015839 2.745797
```

(d)

First center the matrix:

```
X_center <- scale(X, center = TRUE, scale = FALSE)
```

```
X_center_scale <- scale(X, center = T, scale = T)
```

```
( svdObj2 <- svd(X_center) )
```

```
## $d
## [1] 3.8322510 3.3823499 2.0106498 0.4804676
##
## $u
## [,1]      [,2]      [,3]      [,4]
## [1,] -0.5712672 0.1656304 0.28534764 -0.60398428
## [2,] 0.5089398 0.1143288 0.69563542 0.20976309
## [3,] -0.4996900 -0.2490239 -0.08995577 0.69296822
## [4,] 0.3402207 -0.6851706 -0.32383532 -0.33154457
## [5,] 0.2217967 0.6542353 -0.56719197 0.03279753
##
## $v
## [,1]      [,2]      [,3]      [,4]
## [1,] 0.3436766 -0.80716469 0.4461250 0.177042475
## [2,] -0.3682374 0.17791687 0.2582385 0.875248373
## [3,] -0.5751820 0.03345965 0.6822505 -0.450089261
## [4,] 0.6445566 0.56188184 0.5184782 0.003988055
```

```
( svdObj3 <- svd(X_center_scale) )
```

```
## $d
## [1] 3.0431883 2.1387235 1.3876419 0.4892004
##
## $u
## [,1]      [,2]      [,3]      [,4]
## [1,] -0.5985980 -0.003516761 0.1400083 -0.64966586
## [2,] 0.3012174 0.224786627 0.7869330 0.19868439
## [3,] -0.4397787 -0.343527972 -0.1595103 0.68054368
## [4,] 0.5757526 -0.580732932 -0.2401616 -0.27125742
## [5,] 0.1614067 0.702991038 -0.5272694 0.04169521
```

```

## 
## $v
##           [,1]      [,2]      [,3]      [,4]
## [1,]  0.3592298 -0.6917233 0.5565679  0.287584734
## [2,] -0.6339580  0.1305970 0.1900625  0.738185562
## [3,] -0.5983454 -0.1695573 0.4907728 -0.610225552
## [4,]  0.3332176  0.6897200 0.6428456 -0.001368531

```

The principal component directions:

```
svdObj2$v
```

```

##           [,1]      [,2]      [,3]      [,4]
## [1,]  0.3436766 -0.80716469 0.4461250  0.177042475
## [2,] -0.3682374  0.17791687 0.2582385  0.875248373
## [3,] -0.5751820  0.03345965 0.6822505 -0.450089261
## [4,]  0.6445566  0.56188184 0.5184782  0.003988055

```

```
svdObj3$v
```

```

##           [,1]      [,2]      [,3]      [,4]
## [1,]  0.3592298 -0.6917233 0.5565679  0.287584734
## [2,] -0.6339580  0.1305970 0.1900625  0.738185562
## [3,] -0.5983454 -0.1695573 0.4907728 -0.610225552
## [4,]  0.3332176  0.6897200 0.6428456 -0.001368531

```

The “loadings” of the first principal component:

```
svdObj2$v[,1]
```

```
## [1]  0.3436766 -0.3682374 -0.5751820  0.6445566
```

```
svdObj3$v[,1]
```

```
## [1]  0.3592298 -0.6339580 -0.5983454  0.3332176
```

The principal components:

```
X_center %*% svdObj2$v
```

```

##           [,1]      [,2]      [,3]      [,4]
## [1,] -2.1892395  0.5602201  0.5737342 -0.29019489
## [2,]  1.9503851  0.3867000  1.3986792  0.10078437
## [3,] -1.9149375 -0.8422861 -0.1808695  0.33294879
## [4,]  1.3038112 -2.3174867 -0.6511194 -0.15929643
## [5,]  0.8499806  2.2128526 -1.1404244  0.01575815

```

```
X_center_scale %*% svdObj3$v
```

```
##          [,1]          [,2]          [,3]          [,4]
## [1,] -1.8216464 -0.007521379  0.1942814 -0.31781681
## [2,]  0.9166614  0.480756438  1.0919812  0.09719649
## [3,] -1.3383294 -0.734711340 -0.2213431  0.33292225
## [4,]  1.7521236 -1.242027160 -0.3332583 -0.13269924
## [5,]  0.4911909  1.503503442 -0.7316611  0.02039731
```

(e)

The rank = 1 approximation:

```
( rank1approx_center <- best_approx(X_center, 1) )
```

```
##          [,1]          [,2]          [,3]          [,4]
## [1,] -0.7523905  0.8061597  1.2592112 -1.4110888
## [2,]  0.6703018 -0.7182046 -1.1218265  1.2571337
## [3,] -0.6581193  0.7051515  1.1014376 -1.2342856
## [4,]  0.4480894 -0.4801120 -0.7499288  0.8403802
## [5,]  0.2921185 -0.3129946 -0.4888936  0.5478606
```

```
( rank1approx_center_scale <- best_approx(X_center_scale, 1) )
```

```
##          [,1]          [,2]          [,3]          [,4]
## [1,] -0.6543898  1.1548473  1.0899737 -0.6070047
## [2,]  0.3292921 -0.5811248 -0.5484801  0.3054477
## [3,] -0.4807678  0.8484446  0.8007832 -0.4459550
## [4,]  0.6294151 -1.1107727 -1.0483750  0.5838385
## [5,]  0.1764504 -0.3113944 -0.2939018  0.1636735
```

The rank = 2 approximation:

```
( rank1approx_center <- best_approx(X_center, 2) )
```

```
##          [,1]          [,2]          [,3]          [,4]
## [1,] -1.20458032  0.9058323  1.2779560 -1.0963113
## [2,]  0.35817116 -0.6494042 -1.1088876  1.4744134
## [3,]  0.02174434  0.5552946  1.0732550 -1.7075509
## [4,]  2.31868287 -0.8924320 -0.8274711 -0.4617735
## [5,] -1.49401805  0.0807092 -0.4148523  1.7912223
```

```
( rank1approx_center_scale <- best_approx(X_center_scale, 2) )
```

```
##          [,1]          [,2]          [,3]          [,4]
## [1,] -0.649187041  1.1538651  1.0912490 -0.6121924
## [2,] -0.003258318 -0.5183395 -0.6299958  0.6370351
## [3,]  0.027449121  0.7524935  0.9253588 -0.9527000
## [4,]  1.488554221 -1.2729777 -0.8377803 -0.2728124
## [5,] -0.863557983 -0.1150414 -0.5488317  1.2006698
```

(f)

```
cov_matrix_center <- 1/5 * t(X_center) %*% X_center  
cov_matrix_center_scale <- 1/5 * t(X_center_scale) %*% X_center_scale
```

```
( eigenX <- eigen( cov_matrix_center ) )
```

```
## eigen() decomposition  
## $values
```

```
## [1] 2.93722954 2.28805812 0.80854251 0.04616983
```

```
##
```

```
## $vectors
```

```
## [,1]      [,2]      [,3]      [,4]  
## [1,] 0.3436766 0.80716469 -0.4461250 0.177042475  
## [2,] -0.3682374 -0.17791687 -0.2582385 0.875248373  
## [3,] -0.5751820 -0.03345965 -0.6822505 -0.450089261  
## [4,] 0.6445566 -0.56188184 -0.5184782 0.003988055
```

```
( best_approx(cov_matrix_center, 1) )
```

```
## [,1]      [,2]      [,3]      [,4]
```

```
## [1,] 0.3469268 -0.3717198 -0.5806216 0.6506523
```

```
## [2,] -0.3717198 0.3982846 0.6221155 -0.6971509
```

```
## [3,] -0.5806216 0.6221155 0.9717365 -1.0889408
```

```
## [4,] 0.6506523 -0.6971509 -1.0889408 1.2202816
```

```
( best_approx(cov_matrix_center, 2) )
```

```
## [,1]      [,2]      [,3]      [,4]
```

```
## [1,] 1.8376307 -0.7003038 -0.6424162 -0.3870534
```

```
## [2,] -0.7003038 0.4707118 0.6357364 -0.4684178
```

```
## [3,] -0.6424162 0.6357364 0.9742981 -1.0459245
```

```
## [4,] -0.3870534 -0.4684178 -1.0459245 1.9426472
```

```
( eigenX <- eigen( cov_matrix_center_scale ) )
```

```
## eigen() decomposition
```

```
## $values
```

```
## [1] 1.85219898 0.91482763 0.38510998 0.04786341
```

```
##
```

```
## $vectors
```

```
## [,1]      [,2]      [,3]      [,4]
```

```
## [1,] 0.3592298 -0.6917233 -0.5565679 0.287584734
```

```
## [2,] -0.6339580 0.1305970 -0.1900625 0.738185562
```

```
## [3,] -0.5983454 -0.1695573 -0.4907728 -0.610225552
```

```
## [4,] 0.3332176 0.6897200 -0.6428456 -0.001368531
```

```
( best_approx(cov_matrix_center_scale, 1) )
```

```

##          [,1]      [,2]      [,3]      [,4]
## [1,]  0.2390190 -0.4218135 -0.3981182  0.2217114
## [2,] -0.4218135  0.7444038  0.7025869 -0.3912696
## [3,] -0.3981182  0.7025869  0.6631191 -0.3692900
## [4,]  0.2217114 -0.3912696 -0.3692900  0.2056571

```

```
( best_approx(cov_matrix_center_scale, 2) )
```

```

##          [,1]      [,2]      [,3]      [,4]
## [1,]  0.6767468 -0.5044563 -0.2908210 -0.2147486
## [2,] -0.5044563  0.7600067  0.6823293 -0.3088662
## [3,] -0.2908210  0.6823293  0.6894201 -0.4762764
## [4,] -0.2147486 -0.3088662 -0.4762764  0.6408530

```

Problem 2

Making up a matrix:

```

mother_wavelet <- function(x){

  value = ( x > 0 && x <= 1/2 ) - ( x > 1/2 && x <= 1)

  return(value)
}

haar_basis <- function(x){
  index <- 1;
  row <- numeric(16)
  for(m in 0:3)
    for( j in 0:( 2^m - 1 ) ){
      row[index] = sqrt(2^m) * mother_wavelet( (2^m) * ( x - j / (2^m) ) )
      index = index + 1
    }

  return(row)
}

x_trans <- ( data2$x - min(data2$x) )/( max(data2$x) - min(data2$x) )
X_h <- t( apply(as.matrix(x_trans), 1, haar_basis) )

```

(a)

```
( fit_ols <- lm(y~X_h, data = data2) )
```

```

##
## Call:
## lm(formula = y ~ X_h, data = data2)
##
## Coefficients:

```

```

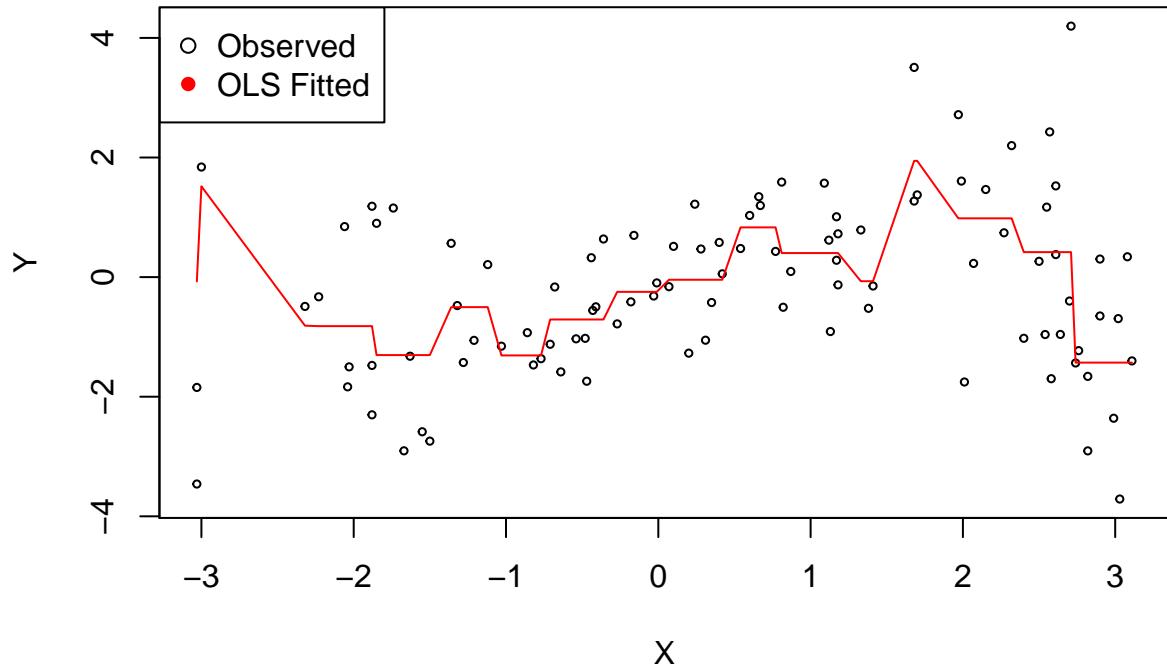
## (Intercept)          X_h1          X_h2          X_h3          X_h4
## -0.07184      -0.45131      0.11927     -0.07014      0.35370
##          X_h5          X_h6          X_h7          X_h8          X_h9
## -0.10709      0.05676      0.49231      0.41207      0.08564
##          X_h10         X_h11         X_h12         X_h13         X_h14
##  0.14280     -0.08164     -0.15481      0.08319      0.16989
##          X_h15         X_h16
##  0.32665           NA

```

```

plot(x = data2$x, y = data2$y, ylab = "Y", xlab = "X",
      pch = 1, cex = 0.5)
lines(x = sort( data2$x ), y = fit_ols$fitted.values[order(data2$x)], pch = 16,
      col = "red", cex = 0.5)
legend("topleft", legend = c("Observed", "OLS Fitted"), col = c("black", "red"),
      pch = c(1, 16))

```



(b)

```

y_centered <- data2$y - mean( data2$y )

get_lambda <- function(M)
{
  fit_lasso <- glmnet(X_h, y_centered, standardize = T, pmax = M, nlambda = 20000)

```

```

    return( min(fit_lasso$lambda) )
}

get_predicted <- function(lambda)
{
  fit_lasso <- glmnet(X_h, y_centered, standardize = T, lambda = lambda)
  value = predict(fit_lasso, X_h)

  return(value)
}

get_coefficient_vector <- function(lambda){
  fit_lasso <- glmnet(X_h, y_centered, standardize = T, lambda = lambda)
  return(fit_lasso$beta)
}

M <- c(2,4,8)
lambda <- apply(as.matrix(M), 1, get_lambda)
value <- apply(as.matrix(lambda), 1, get_predicted) + mean(data2$y)

( coef <- apply( as.matrix(lambda), 1, get_coefficient_vector) )

## [[1]]
## 16 x 1 sparse Matrix of class "dgCMatrix"
##           s0
## V1   -0.03994013
## V2   .
## V3   .
## V4   .
## V5   .
## V6   .
## V7   .
## V8   .
## V9   .
## V10  .
## V11  .
## V12  .
## V13  .
## V14  .
## V15  0.02880334
## V16  .
##
## [[2]]
## 16 x 1 sparse Matrix of class "dgCMatrix"
##           s0
## V1   -0.30242980
## V2   .
## V3   .
## V4   0.09544829
## V5   .
## V6   .
## V7   0.25349740
## V8   .

```

```

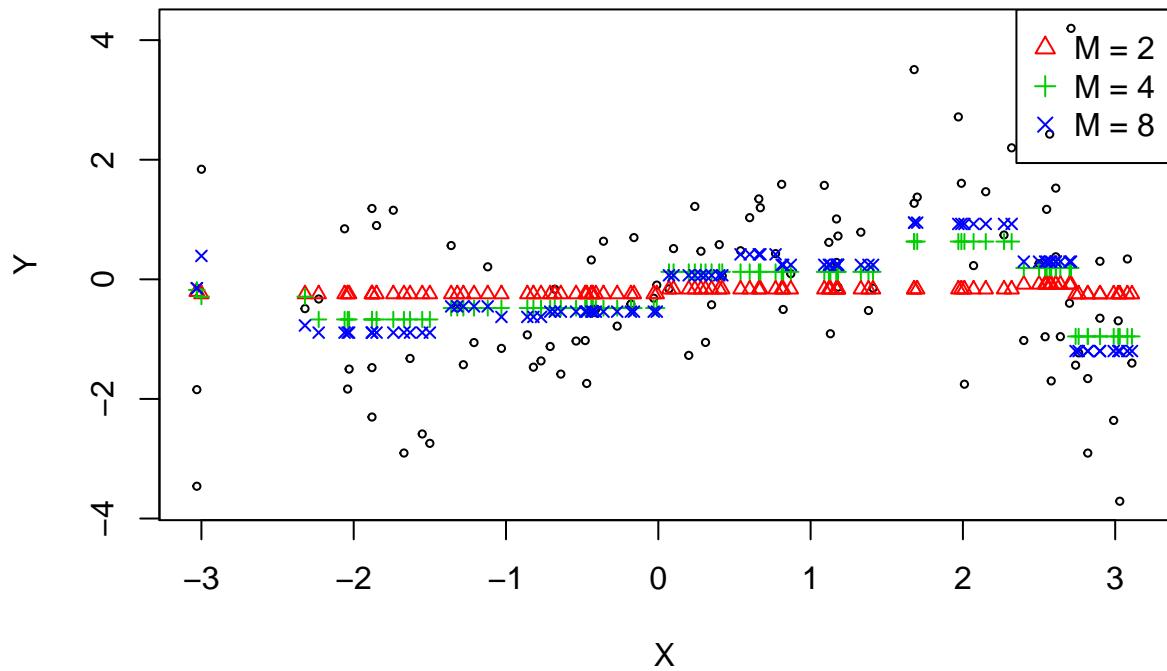
## V9   .
## V10  .
## V11  .
## V12  .
## V13  .
## V14  .
## V15  0.20237856
## V16  .
##
## [[3]]
## 16 x 1 sparse Matrix of class "dgCMatrix"
##           s0
## V1  -0.391467689
## V2  .
## V3  .
## V4  0.175157628
## V5  .
## V6  .
## V7  0.347304902
## V8  0.205501442
## V9  .
## V10 0.031114449
## V11  .
## V12 -0.061899306
## V13  .
## V14 0.003836725
## V15 0.264369490
## V16  .

plot(x = data2$x, y = data2$y, ylab = "Y", xlab = "X",
      pch = 1, cex = 0.5)

points(x = data2$x, y = value[,1], cex = 0.75, pch = 2, col = 2)
points(x = data2$x, y = value[,2], cex = 0.75, pch = 3, col = 3)
points(x = data2$x, y = value[,3], cex = 0.75, pch = 4, col = 4)

legend("topright", legend = c("M = 2", "M = 4", "M = 8"), col = c(2,3,4), pch = c(2,3,4))

```



Problem 3

```

knot <- c(0, .1, .3, .5, .7, .9, 1)
K <- length(knot)

only_positive <- function(x){
  value <- ifelse(x >= 0, x, 0)

  return(value)
}

make_row <- function(x){
  row <- numeric(K)
  row[1] = 1
  row[2] = x

  for(j in 1:(K - 2)){
    value = only_positive( ( x - knot[j] )^3 )
    - ( knot[K] - knot[j] ) / ( knot[K] - knot[K-1] ) * only_positive( (x - knot[K-1])^3 )
    + ( knot[K-1] - knot[j] ) / ( knot[K] - knot[K-1] ) * only_positive( (x - knot[K])^3 )

    row[j+2] = value
  }
}

```

```

    return(row)
}

x_trans <- ( data2$x - min(data2$x) )/( max(data2$x) - min(data2$x) )
X_h_3 <- t( apply(as.matrix(x_trans), 1, make_row) )

```

Then to find $\hat{\beta}^{OLS}$:

```

fit_ols <- lm(data2$y ~ 0 + X_h_3)
fit_ols$coefficients

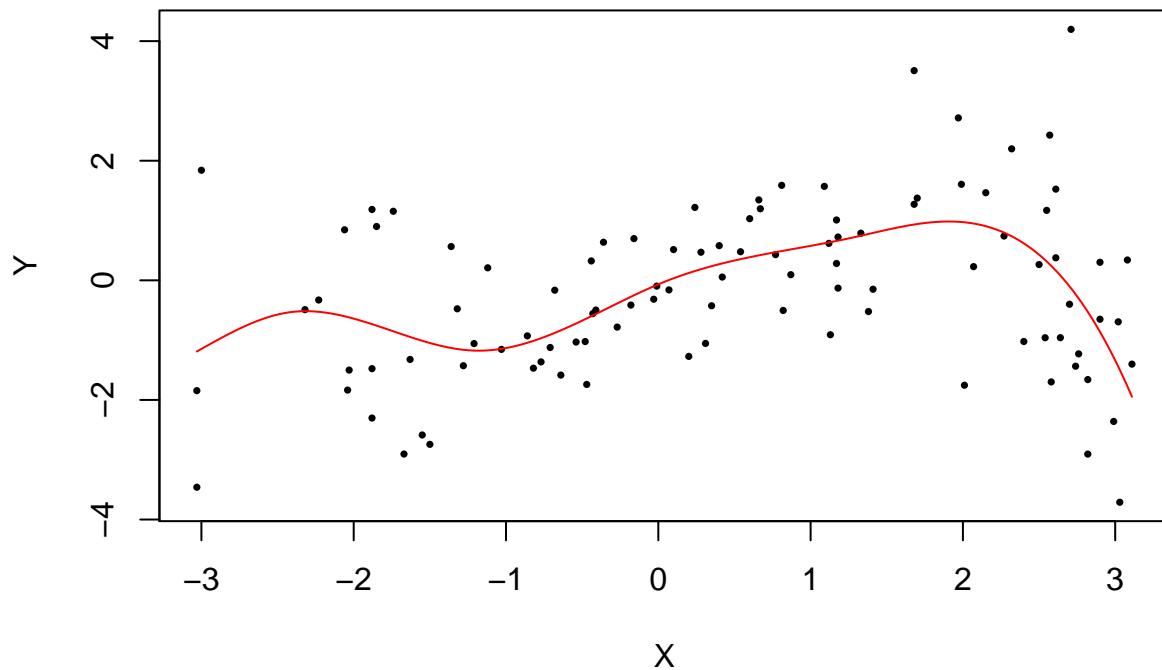
##      X_h_31      X_h_32      X_h_33      X_h_34      X_h_35      X_h_36
## -1.190293     8.843849 -226.408590   434.236477 -343.350843  193.448300
##      X_h_37
## -228.230208

x_plot <- seq(from = min( data2$x ), to = max( data2$x ), length.out = 500)
x_plot_trans <- ( x_plot - min(data2$x) )/( max(data2$x) - min(data2$x) )
x_plot_h <- t( apply(as.matrix(x_plot_trans), 1, make_row) )
y_plot <- x_plot_h %*% fit_ols$coefficients

plot(x = data2$x, y = data2$y, pch = 16, cex = 0.5,
      ylab = "Y", xlab = "X", main = "Natural Cubic Regression Spline")
lines(x = x_plot, y = y_plot, col = "red")

```

Natural Cubic Regression Spline



Problem 4

(a)

```
y <- c(0, 1.5, 2, 0.5, 0, -0.5,
      0, 1.5, 3.5, 4.5, 3.5)

Omiga <- matrix(0, nrow = N, ncol = N)

for( j in 1:(N-2) )
  for( k in j:(N-2)){
    if(k == j){
      Omiga[j+2, k+2] = 12 * (x[N-1] - x[j])^2 * (x[N] - x[j])
    } else {
      Omiga[j+2, k+2] = 6 * ( x[N-1] - x[k] )^2 * ( 2*x[N-1] + x[k] - 3*x[j] ) +
        12 * ( x[N-1] - x[k] ) * ( x[N-1] - x[j] ) * ( x[N] - x[N-1] )
      Omiga[k+2, j+2] = Omiga[j+2, k+2]
    }
  }

H_matrix <- matrix(nrow = N, ncol = N)
for(i in 1:N)
  for(j in 1:N){
    H_matrix[i, j] <- compute_H(i, j)
  }
H_matrix

##      [,1]  [,2]  [,3]  [,4]  [,5]  [,6]  [,7]  [,8]  [,9]  [,10] [,11]
## [1,]     1  0.0  0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000
## [2,]     1  0.1  0.001 0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000
## [3,]     1  0.2  0.008 0.001 0.000 0.000 0.000 0.000 0.000 0.000 0.000
## [4,]     1  0.3  0.027 0.008 0.001 0.000 0.000 0.000 0.000 0.000 0.000
## [5,]     1  0.4  0.064 0.027 0.008 0.001 0.000 0.000 0.000 0.000 0.000
## [6,]     1  0.5  0.125 0.064 0.027 0.008 0.001 0.000 0.000 0.000 0.000
## [7,]     1  0.6  0.216 0.125 0.064 0.027 0.008 0.001 0.000 0.000 0.000
## [8,]     1  0.7  0.343 0.216 0.125 0.064 0.027 0.008 0.001 0.000 0.000
## [9,]     1  0.8  0.512 0.343 0.216 0.125 0.064 0.027 0.008 0.001 0.000
## [10,]    1  0.9  0.729 0.512 0.343 0.216 0.125 0.064 0.027 0.008 0.001
## [11,]    1  1.0  0.990 0.720 0.504 0.336 0.210 0.120 0.060 0.024 0.006

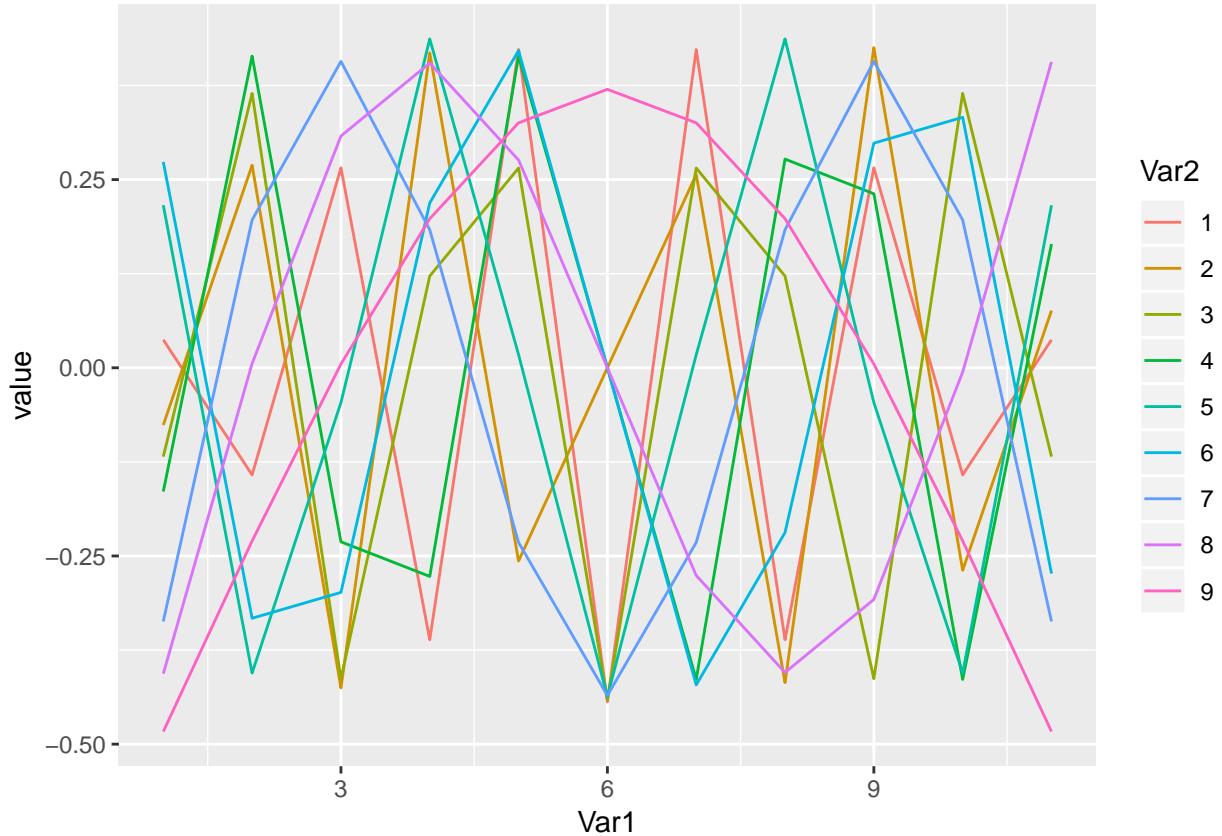
K_matrix <- solve( t(H_matrix) ) %*% Omiga %*% solve( H_matrix )
K_matrix <- ( K_matrix + t(K_matrix) )/2
```

(b)

```
K_eigen <- eigen(K_matrix); K_eigen$values

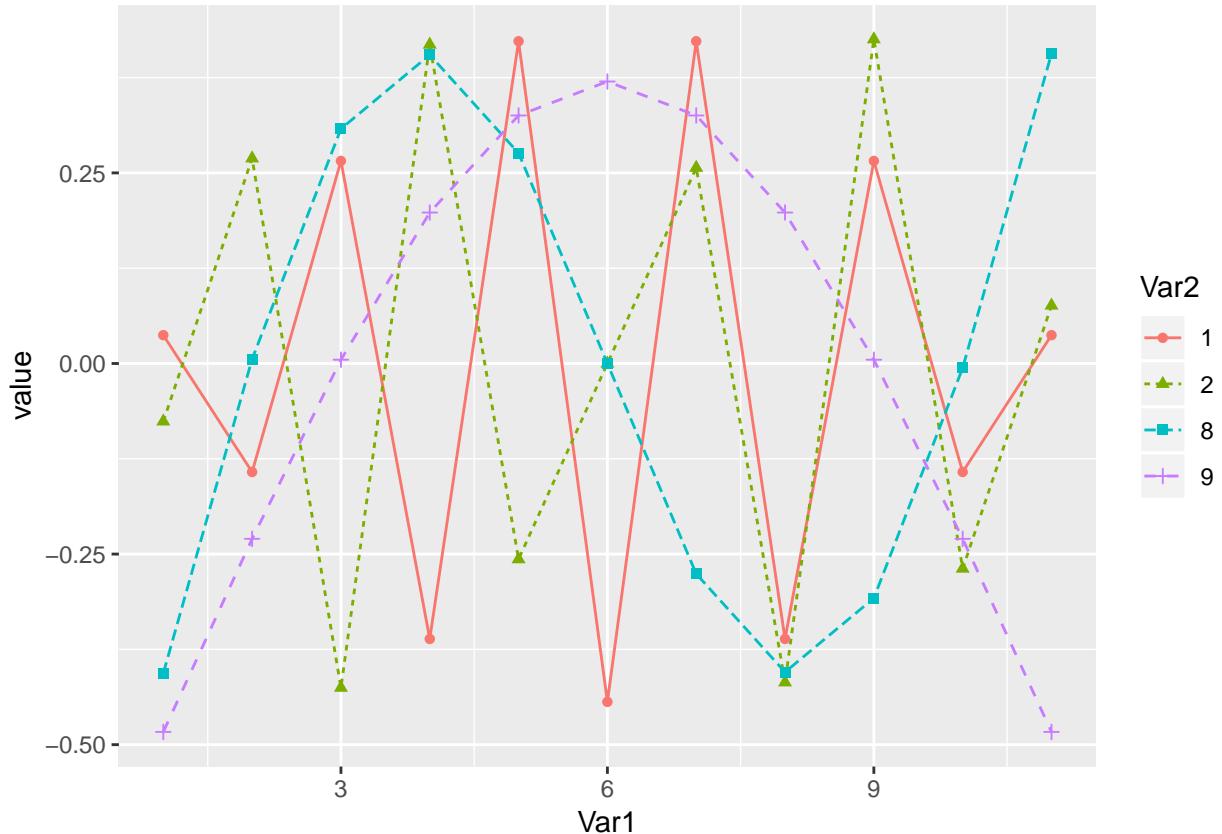
## [1] 4.366292e+04 3.333961e+04 2.200669e+04 1.282846e+04 6.610687e+03
## [6] 2.938555e+03 1.060305e+03 2.713642e+02 3.508288e+01 -1.958878e-12
## [11] -3.244671e-12
```

```
reshape2::melt(K_eigen$vectors) %>% mutate(Var2 = as.factor(Var2)) %>%
  filter(Var2 %in% 1:9) %>% ggplot() + geom_line(aes(x = Var1, y = value, col = Var2))
```



A large eigen value of K corresponds to a small eigen value of S_λ . From the plot below, we can draw the conclusion that the $Y[c(2,3,6,9)]$ will get suppressed most in the smoothing.

```
reshape2::melt(K_eigen$vectors) %>%
  mutate(Var2 = as.factor(Var2)) %>%
  filter(Var2 %in% c(1:2, 8:9)) %>%
  ggplot(aes(x = Var1, y = value, col = Var2)) + geom_line(aes(linetype = Var2)) + geom_point(aes(pch
```



(c)

```

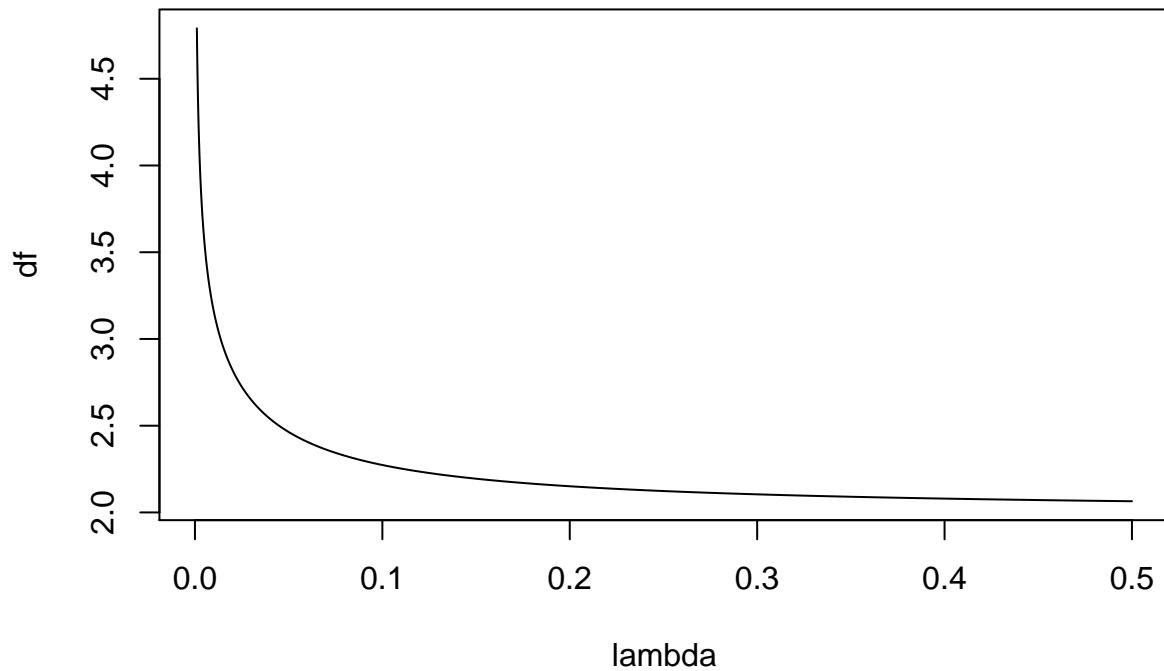
compute_df <- function(lambda, eta_array, target = 0){
  eta_array[10:11] = 0
  value <- sum( 1/(1 + lambda * eta_array) )

  return( value-target )
}

lambda <- as.matrix( seq(from = 0.001, to = 0.5, by = 0.0001) )
df <- apply(lambda, 1, compute_df, eta_array = K_eigen$values)

plot(x = lambda, y = df, type = "l")

```



```

df <- c(2.5, 3, 4, 5)
lambda_roots_4 <- numeric(length(df))

for(i in 1:length(df)){
  lambda_roots_4[i] <- uniroot(function(x) compute_df(x, eta_array = K_eigen$values, target = df[i]), i)
}
lambda_roots_4

## [1] 0.0448424394 0.0137079634 0.0026479555 0.0008216407

```

Problem 5

(a)

```

x <- seq(0,1,by=0.1)
N <- length(x)
B <- as.matrix( cbind(1,x) )

compute_K_lambda <- function(lambda, x0, xi){
  value <- dnorm( (x0 - xi)/lambda )

  return(value)
}

```

```

}

compute_W <- function(x0, lambda){
  value <- numeric(N)
  for(i in 1:N){
    value[i] <- compute_K_lambda(lambda, x0, x[i])
  }

  return(diag(value))
}

compute_I <- function(lambda, x0){
  W <- compute_W(x0, lambda)

  part1 <- matrix(c(1,x0), nrow = 1)
  part2 <- solve( t(B) %*% W %*% B)
  part3 <- t(B) %*% W

  value <- part1 %*% part2 %*% part3

  return(value)
}

get_trace_L <- function(lambda, target = 0){
  L <- apply(as.matrix( x ), 1, compute_I, lambda = lambda)

  return( tr(L) - target )
}

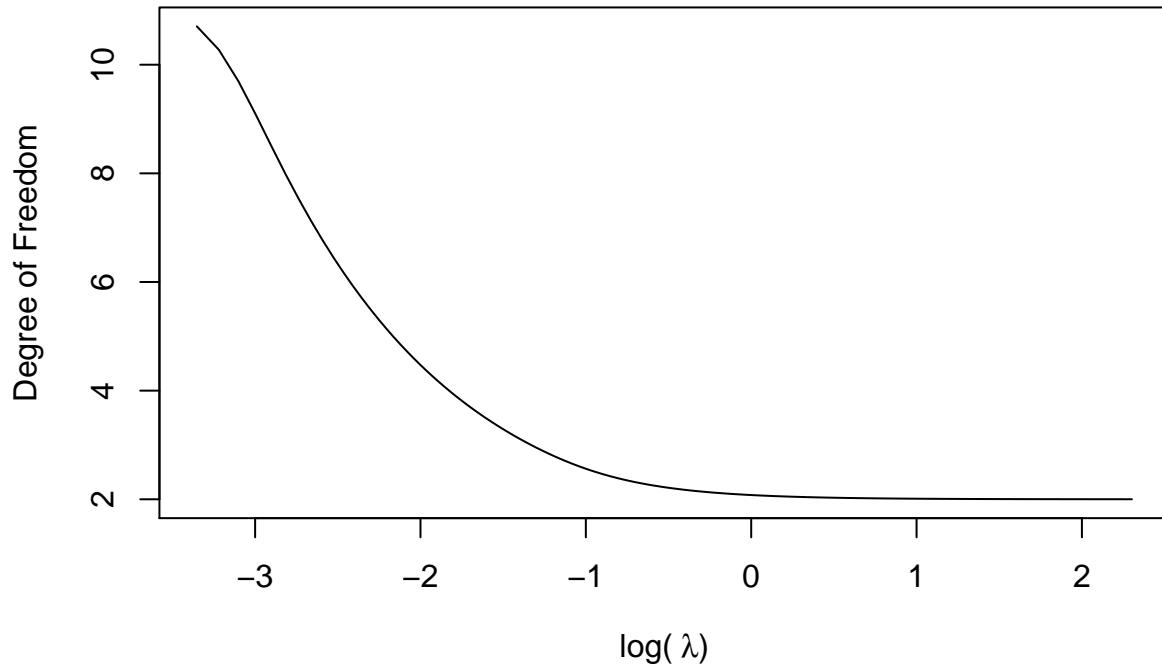
```

This is the plot of effective degree of freedom v.s. $\log(\lambda)$:

```

lambda <- seq(0.035,10,length.out = 2000)
trace_I <- apply(as.matrix( lambda ), 1,get_trace_L)
plot(log(lambda), trace_I, ty = "l", ylab = "Degree of Freedom",
     xlab = expression( paste( "log(", ~lambda, ")" ) )
)

```



Find out the value of alpha by solving the root:

```
df <- c(2.5, 3, 4, 5)
lambda_roots <- numeric(length(df))

for(i in 1:length(df)){
  lambda_roots[i] <- uniroot(function(x) get_trace_L(x, target = df[i]), interval = c(0.05,4))$root
}

lambda_roots
```

[1] 0.3915776 0.2642124 0.1609244 0.1146486

(b)

```
lambda5 <- lambda_roots[3]
lambda4 <- lambda_roots_4[3]

L_lambda5 <- t( apply(as.matrix( x ), 1, compute_I, lambda = lambda5) )
S_lambda4 <- solve( diag(1, nrow = 11) + lambda4*K_matrix )
```

The matrix difference is:

```
L_lambda5 - S_lambda4
```

```
##          [,1]          [,2]          [,3]          [,4]          [,5]
## [1,]  0.090027536 -0.011577964 -0.094409573 -0.053444610  0.009269582
## [2,]  0.028821277  0.023347628 -0.019283615 -0.042654653 -0.026454779
## [3,]  0.038764026  0.015740216 -0.020661092 -0.028294980 -0.027429832
## [4,]  0.052673668  0.007797074 -0.015630840 -0.036983145 -0.025077386
## [5,]  0.052563434  0.007659622 -0.011710767 -0.021773047 -0.035424474
## [6,]  0.040489059  0.013278246 -0.005564228 -0.012537779 -0.019086010
## [7,]  0.024457299  0.016186744  0.005322851 -0.006004389 -0.011640453
## [8,]  0.011086811  0.013664341  0.012591005  0.004407607 -0.008059208
## [9,]  0.002671861  0.008414849  0.012715407  0.011936735  0.001739538
## [10,] -0.001816245  0.003047093  0.008382383  0.013235838  0.013155994
## [11,] -0.004492881 -0.001816568  0.002662322  0.010931083  0.022859007
##          [,6]          [,7]          [,8]          [,9]          [,10]
## [1,]  0.0299920680  0.022859007  0.010931083  0.002662322 -0.001816568
## [2,]  0.0002190787  0.013155994  0.013235838  0.008382383  0.003047093
## [3,] -0.0155967283  0.001739538  0.011936735  0.012715407  0.008414849
## [4,] -0.0164699248 -0.008059208  0.004407607  0.012591005  0.013664341
## [5,] -0.0196368197 -0.011640453 -0.006004389  0.005322851  0.016186744
## [6,] -0.0331585750 -0.019086010 -0.012537779 -0.005564228  0.013278246
## [7,] -0.0196368197 -0.035424474 -0.021773047 -0.011710767  0.007659622
## [8,] -0.0164699248 -0.025077386 -0.036983145 -0.015630840  0.007797074
## [9,] -0.0155967283 -0.027429832 -0.028294980 -0.020661092  0.015740216
## [10,]  0.0002190787 -0.026454779 -0.042654653 -0.019283615  0.023347628
## [11,]  0.0299920680  0.009269582 -0.053444610 -0.094409573 -0.011577964
##          [,11]
## [1,] -0.004492881
## [2,] -0.001816245
## [3,]  0.002671861
## [4,]  0.011086811
## [5,]  0.024457299
## [6,]  0.040489059
## [7,]  0.052563434
## [8,]  0.052673668
## [9,]  0.038764026
## [10,]  0.028821277
## [11,]  0.090027536
```

The plots are:

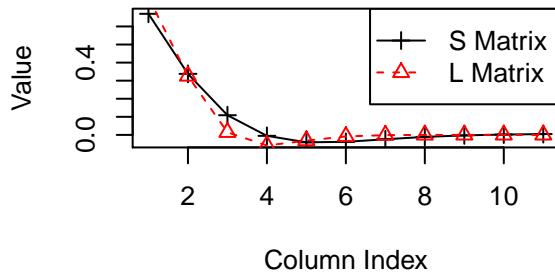
```
par(mfrow = c(2,2))

for(i in c(1,3,5)){
  plot(x = 1:11, y = S_lambda4[i, ], xlab = "Column Index", col = 1,
    main = paste("Row", i, " of the Smoothing Matrix"), pch = 3, ylab = "Value")
  lines(x = 1:11, y = S_lambda4[i, ], col = 1, lty = 1)

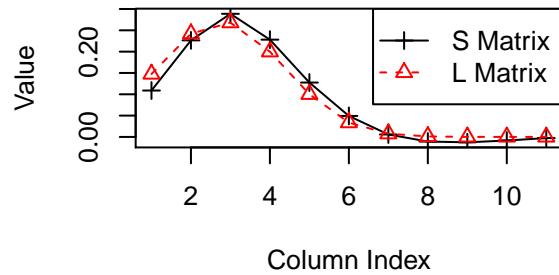
  points(x = 1:11, y = L_lambda5[i, ], pch = 2, col = 2)
  lines(x = 1:11, y = L_lambda5[i, ], col = 2, lty = 2)

  legend("topright", legend = c("S Matrix", "L Matrix"), pch = c(3, 2), col = c(1, 2), lty = c(1, 2))
}
```

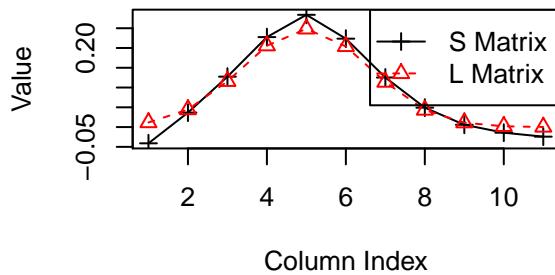
Row 1 of the Smoothing Matrix



Row 3 of the Smoothing Matrix



Row 5 of the Smoothing Matrix



From the plots above, the two smoothing matrices are close to each other when the effective degree of freedom is the same.

6

(a)

The largest effective degrees of freedom I can get for the locally weighted linear regression smoother with tricube kernel is about 4.72. If I try to go over it, the effective degrees of freedom will jump to 11, which means the fitted values are the same as the observed values.

```
lambda_df5 <- uniroot(function(lambda) smooth.spline(x, y, lambda = lambda)$df-5,
                        interval = c(0.0001, 1))$root
lambda_df9 <- uniroot(function(lambda) smooth.spline(x, y, lambda = lambda)$df-9,
                        interval = c(0.000001, 1))$root
cubic_smoothing_5 <- smooth.spline(x, y, lambda = lambda_df5)
cubic_smoothing_9 <- smooth.spline(x, y, lambda = lambda_df9)

tricube_5 <- loess(y~x, degree = 1, enp.target = 5)
( tricube_9 <- loess(y~x, degree = 1, enp.target = 9) )
```

```
## Call:
## loess(formula = y ~ x, enp.target = 9, degree = 1)
##
```

```

## Number of Observations: 11
## Equivalent Number of Parameters: 11
## Residual Standard Error: NaN

```

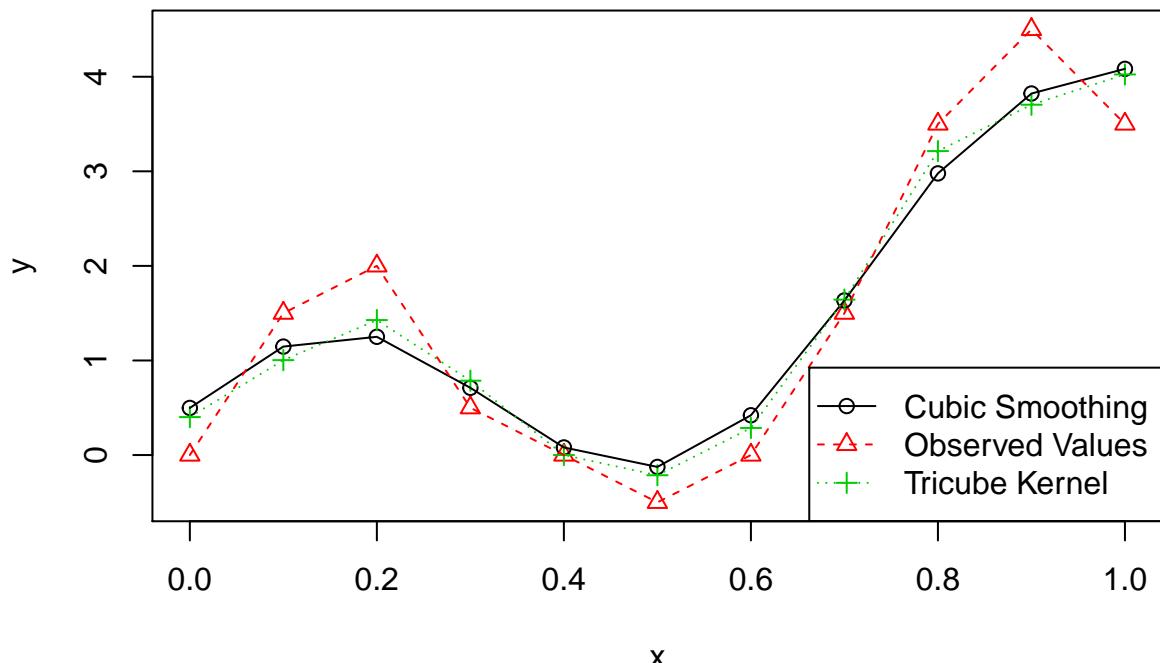
This is the plot for approximately 5 effective degrees of freedom:

```

plot(x, y, pch = 2, col = 2)
points(x, cubic_smoothing_5$y, ylab = "Y", xlab = "x")
points(x, tricube_5$fitted, pch = 3, col = 3)

lines(x, cubic_smoothing_5$y)
lines(x, tricube_5$fitted, lty = 3, col = 3)
lines(x, y, lty = 2, col = 2)
legend("bottomright", legend = c("Cubic Smoothing", "Observed Values", "Tricube Kernel"),
       pch = c(1, 2, 3), col = c(1, 2, 3), lty = c(1, 2, 3))

```



This is the plot for approximately 9 effective degrees of freedom:

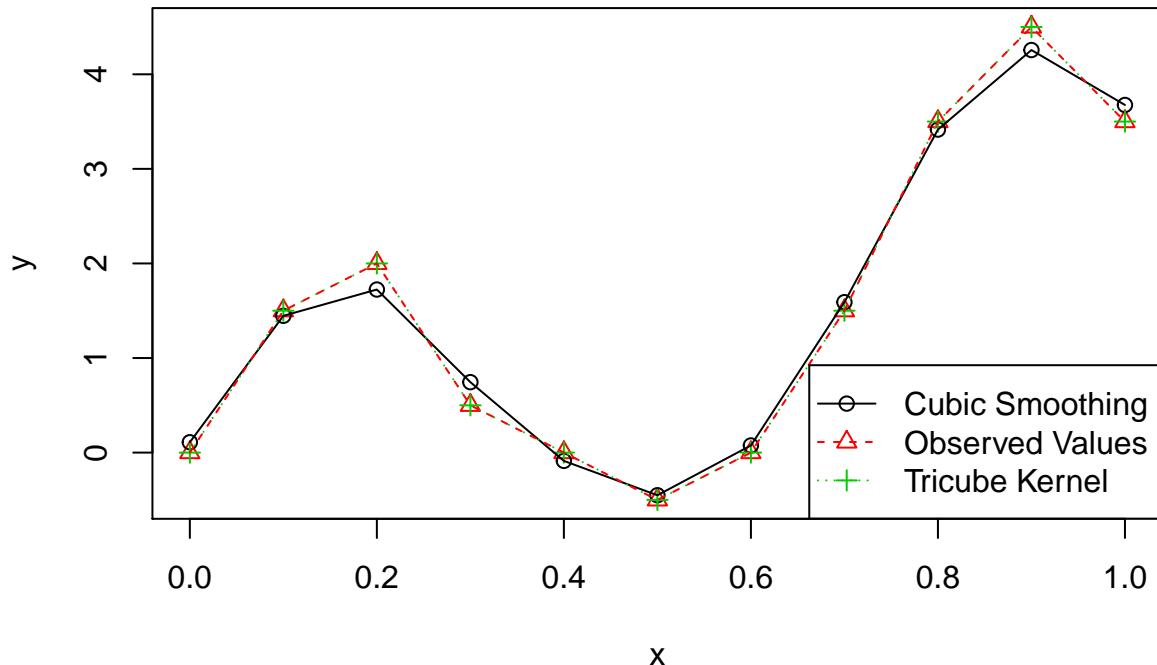
```

plot(x, y, pch = 2, col = 2)
points(x, cubic_smoothing_9$y, ylab = "Y", xlab = "x")
points(x, tricube_9$fitted, pch = 3, col = 3)

lines(x, cubic_smoothing_9$y)
lines(x, tricube_9$fitted, lty = 3, col = 3)
lines(x, y, lty = 2, col = 2)

```

```
legend("bottomright", legend = c("Cubic Smoothing", "Observed Values", "Tricube Kernel"),
      pch = c(1, 2, 3), col = c(1, 2, 3), lty = c(1, 2, 3))
```



We can see that, the fitted value from locally weighted linear regression is the same as the observed one. The reason for that is explained above. Maybe it is caused by some numerical issues.

Problem 7

```
wineknn<- train(y=wine$quality,
                  x=wine[,-12],
                  method="knn",
                  preProcess=c("center","scale"),
                  trControl=trainControl(method="repeatedcv",repeats=100,number=10))

winenet<- train(y=wine$quality,
                  x=wine[,-12],
                  method="glmnet",
                  preProcess=c("center","scale"),
                  trControl=trainControl(method="repeatedcv",repeats=100,number=10))

winepcr<- train(y=wine$quality,
                  x=wine[,-12],
                  method="pcr",
```

```

preProcess=c("center","scale"),
trControl=trainControl(method="repeatedcv",repeats=100,number=10))

winepls<- train(y=wine$quality,
                  x=wine[,-12],
                  method="pls",
                  preProcess=c("center","scale"),
                  trControl=trainControl(method="repeatedcv",repeats=100,number=10))

winemars<- train(y=wine$quality,
                  x=wine[,-12],
                  method="earth",
                  preProcess=c("center","scale"),
                  trControl=trainControl(method="repeatedcv",repeats=100,number=10))

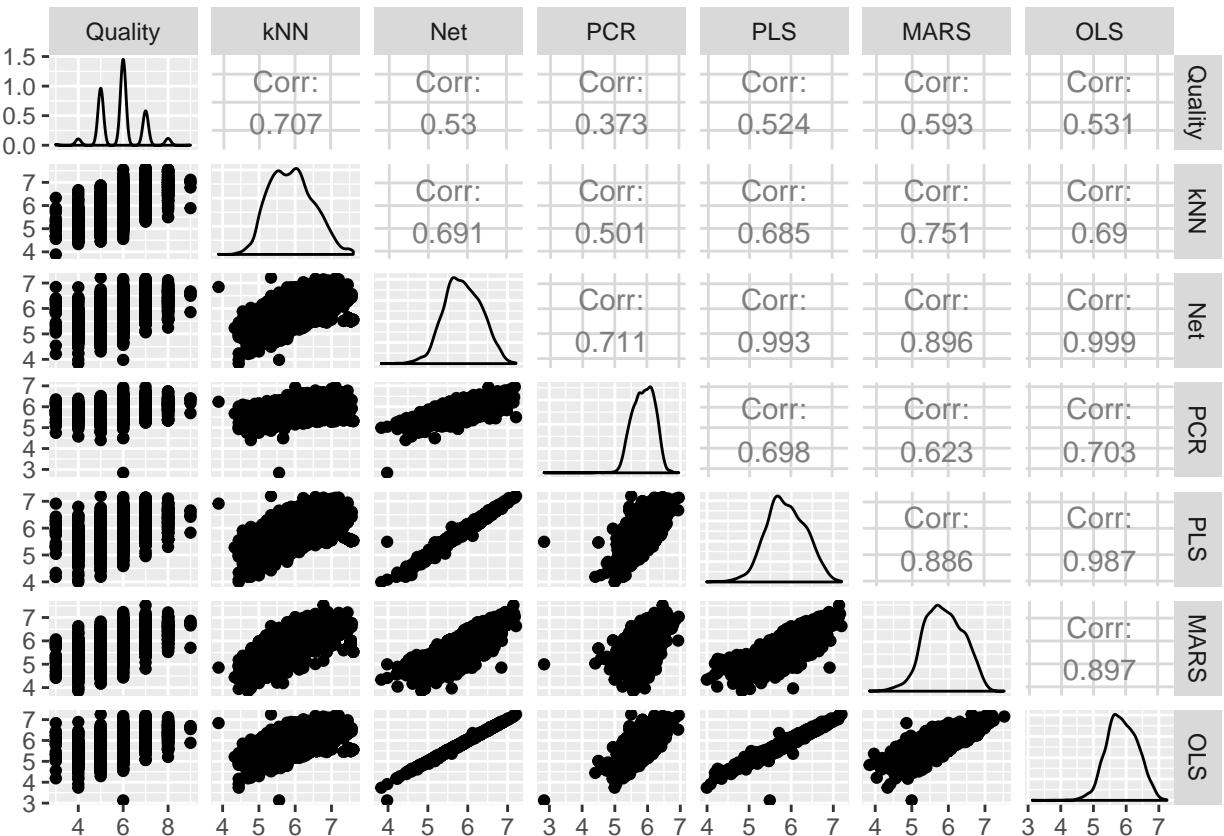
wineols<- lm(quality~.,data=wine)

winepred<- sapply(list(wineknn,winenet,winepcr,winepls,winemars,wineols),predict)
colnames(winepred)<- c("kNN", "Net", "PCR", "PLS", "MARS", "OLS")

res<- data.frame(cbind(Quality=wine$quality,winepred))

ggpairs(res)

```



```

round(cor(res), digits=2)

##          Quality kNN Net PCR PLS MARS OLS
## Quality     1.00 0.71 0.53 0.37 0.52 0.59 0.53
## kNN         0.71 1.00 0.69 0.50 0.69 0.75 0.69
## Net         0.53 0.69 1.00 0.71 0.99 0.90 1.00
## PCR         0.37 0.50 0.71 1.00 0.70 0.62 0.70
## PLS         0.52 0.69 0.99 0.70 1.00 0.89 0.99
## MARS        0.59 0.75 0.90 0.62 0.89 1.00 0.90
## OLS         0.53 0.69 1.00 0.70 0.99 0.90 1.00

```

Problem 8

(a)

```

x <- seq(0, 1, by = 0.1)
y <- c(0, 1.5, 2, 0.5, 0, -0.5, 0, 1.5, 3.5, 4.5, 3.5)

x_std <- scale(x, center = T, scale = T)
y_ctr <- scale(y, center = T, scale = FALSE)

K <- as.matrix( -1/2 * dist(x_std)^2 ) %>% exp

G <- matrix(nrow = 11, ncol = 11)
for(i in 1:11)
  for(j in 1:11){
    G[i, j] = K[i, j] - colMeans(K)[i] - rowMeans(K)[j] + mean(K)
  }
G

##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,]  0.72982684  0.59948554  0.40416853  0.17923959 -0.03591672
## [2,]  0.59948554  0.55801817  0.44007432  0.26283433  0.05920771
## [3,]  0.40416853  0.44007432  0.41100439  0.31113756  0.15519990
## [4,]  0.17923959  0.26283433  0.31113756  0.30014465  0.22158014
## [5,] -0.03591672  0.05920771  0.15519990  0.22158014  0.23188956
## [6,] -0.20963926 -0.13330264 -0.02578076  0.08828844  0.17597102
## [7,] -0.32445509 -0.28406202 -0.19532794 -0.06972905  0.06564248
## [8,] -0.37719216 -0.37623188 -0.32344136 -0.21663027 -0.06972905
## [9,] -0.37505911 -0.40766663 -0.39430890 -0.32344136 -0.19532794
## [10,] -0.33090034 -0.38745656 -0.40766663 -0.37623188 -0.28406202
## [11,] -0.25955782 -0.33090034 -0.37505911 -0.37719216 -0.32445509
##           [,6]      [,7]      [,8]      [,9]      [,10]
## [1,] -0.20963926 -0.32445509 -0.37719216 -0.37505911 -0.33090034
## [2,] -0.13330264 -0.28406202 -0.37623188 -0.40766663 -0.38745656
## [3,] -0.02578076 -0.19532794 -0.32344136 -0.39430890 -0.40766663
## [4,]  0.08828844 -0.06972905 -0.21663027 -0.32344136 -0.37623188
## [5,]  0.17597102  0.06564248 -0.06972905 -0.19532794 -0.28406202
## [6,]  0.20892640  0.17597102  0.08828844 -0.02578076 -0.13330264
## [7,]  0.17597102  0.23188956  0.22158014  0.15519990  0.05920771

```

```

## [8,]  0.08828844  0.22158014  0.30014465  0.31113756  0.26283433
## [9,] -0.02578076  0.15519990  0.31113756  0.41100439  0.44007432
## [10,] -0.13330264  0.05920771  0.26283433  0.44007432  0.55801817
## [11,] -0.20963926 -0.03591672  0.17923959  0.40416853  0.59948554
##          [,11]
## [1,] -0.25955782
## [2,] -0.33090034
## [3,] -0.37505911
## [4,] -0.37719216
## [5,] -0.32445509
## [6,] -0.20963926
## [7,] -0.03591672
## [8,]  0.17923959
## [9,]  0.40416853
## [10,]  0.59948554
## [11,]  0.72982684

```

(c)

```

G_eigen <- eigen(G)

compute_T <- function(z, x){
  return( exp(-1/2 * ( x - z )^2 ) )
}

compute_S <- function(z, i){
  T_xi <- compute_T( z, x_std[i] )

  M <- 0
  for(j in 1:11){
    M <- M + compute_T( z, x_std[j] )
  }
  M <- M/11

  return(T_xi + M)
}

```

The function is becoming closer to 0 with decreasing eigenvalues, but the smallest eigen value is an outlier.

```

X_vector <- NULL
Y_vector <- NULL
Eigen_vector <- NULL

for(i in 1:11){
  eigen_value <- G_eigen$values[i]

  x_plot <- seq(from = min( x_std ), to = max( x_std ), by = 0.01)
  y_plot <- numeric( length(x_plot) )

  Eigen_vector <- c( Eigen_vector, rep(eigen_value, times = length(x_plot)) )
  X_vector <- c(X_vector, x_plot)
}

```

```

for(ii in 1:length(x_plot) ){
  y_value <- 0
  for(jj in 1:11){
    y_value <- y_value + G_eigen$vectors[,i][jj] * compute_S( x_plot[ii], jj )
  }
  y_plot[ii] <- y_value
}

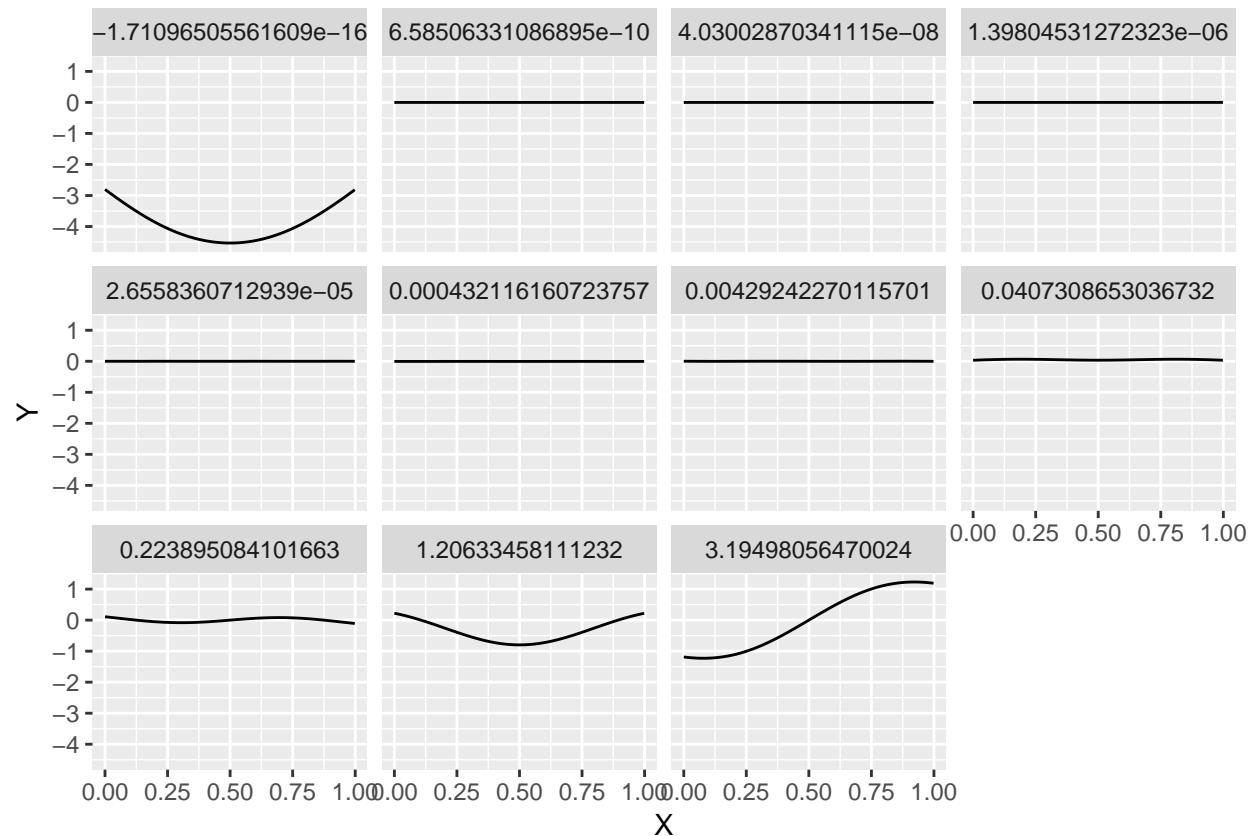
Y_vector <- c(Y_vector, y_plot)
}

X_vector <- ( sd(x) * X_vector + mean(x) )

plot_frame <- data.frame(X = X_vector, Y = Y_vector, Lambda = Eigen_vector)

ggplot(data = plot_frame, mapping = aes(x = X, y = Y) )+ geom_line() + facet_wrap(~Lambda)

```



(d)

The eigen decomposition is equivalent to SVD in this problem. The following is the plot:

```

compute_S065 <- function(y)
{
  value <- compute_T(y, 0.65)
}

```

```

for(i in 1:11){
  value <- value - 1/11 * compute_T( y, x_std[i] )
}

return(value)
}

for(i in 1:length(x_plot))
{
  y_plot[i] <- compute_S065(x_plot[i])
}

kernel_inner <- function(x, y){
  return( exp(-(x - y)^2/2) )
}

S_65_S_xj <- function(j){
  value <- kernel_inner(0.65, x_std[j])

  part2 <- 0
  for(i in 1:11){
    part2 <- part2 + kernel_inner(x_std[i], 0.65) + kernel_inner(x_std[i], x_std[j])
  }
  part2 <- -1/11 * part2

  part3 <- 0
  dict <- expand.grid(x_std, x_std)
  for(i in 1:121){
    part3 <- part3 + kernel_inner(dict[i,1], dict[i,2])
  }
  part3 <- part3 / 121

  return(value + part2 + part3)
}

compute_s_weight <- function(y, j)
{
  weight <- rowSums(G_eigen$vectors)
  value <- weight[j] * S_65_S_xj(j) * compute_S(y, j)
  return(value)
}

final_function <- function(y){
  value = 0;
  for(i in 1:11){
    value = value + compute_s_weight(y, i)
  }
  return(value)
}

y_proj <- numeric(length(x_plot))
for(i in 1:length(x_plot)){
  y_proj <- final_function(x_plot)
}

```

```
}
```

```
data.frame(X = sd(x) * x_plot+mean(x), Original = y_plot, Projection = y_proj) %>% reshape2::melt(id.v
```

